

Master's programme in Computer, Communication and Information Sciences

Kubernetes inter-pod container isolation

Aarni Halinen

© 2023

This work is licensed under a [Creative Commons](#)
“Attribution-NonCommercial-ShareAlike 4.0 International” license.



Author Aarni Halinen

Title Kubernetes inter-pod container isolation

Degree programme Computer, Communication and Information Sciences

Major Computer Science

Supervisor Prof. Mario Di Francesco

Advisors M.Sc. (Tech.) José Luis Martin Navarro, M.Sc. (Tech.) Jacopo Bufalino

Date 1 June 2023

Number of pages 51+9

Language English

Abstract

The abstract is a short description of the essential contents of the thesis: what was studied and how and what were the main findings. For a Finnish thesis, the abstract should be written in both Finnish and English; for a Swedish thesis, the abstract should be written in both Swedish and English. The abstracts for English theses written by Finnish or Swedish speakers should be written in English and either in Finnish or in Swedish, depending on the student's language of basic education. Students educated in languages other than Finnish or Swedish write the abstract only in English. Students may include a second or third abstract in their native language, if they wish. The abstract text of this thesis is written on the readable abstract page and in the pdf file's metadata via the `thesisabstract` macro (see the comment in the TeX file). Write in this the text that goes into the metadata. The metadata cannot contain special characters, linebreak or paragraph break characters, so these must not be used here. If your abstract does not contain special characters and does not require paragraphs, you may take advantage of the `abstracttext` macro (see the comment in the TeX file below). Otherwise, the metadata abstract text must be identical to the text on the abstract page.

Keywords Kubernetes, Container, Docker, Security

Tekijä Aarni Halinen

Työn nimi Opinnäytteen otsikko

Koulutusohjelma Computer, Communication and Information Sciences

Pääaine Computer Science

Työn valvoja Prof. Mario Di Francesco

Työn ohjaajat DI José Luis Martin Navarro, DI Jacopo Bufalino

Päivämäärä 1.6.2023

Sivumäärä 51+9

Kieli englanti

Tiivistelmä

Tiivistelmä on lyhyt kuvaus työn keskeisestä sisällöstä: mitä tutkittiin ja miten sekä mitkä olivat tärkeimmät tulokset. Suomenkielisen opinnäytteen tiivistelmä kirjoitetaan suomeksi ja englanniksi ja ruotsinkielisen vastaavasti ruotsiksi ja englanniksi. Suomen- tai ruotsinkielisten opiskelijoiden, joiden opinnäytteen kieli on englanti, tulee kirjoittaa tiivistelmänsä englanniksi ja koulusivistyskielellään. Muiden kuin koulusivistyskieleltään suomen- tai ruotsinkielisten tulee kirjoittaa tiivistelmänsä vain englanniksi. Opiskelija voi halutessaan lisätä opinnäytteeseensä toisen tai kolmannen tiivistelmän omalla äidinkielellään. Tämän opinnäytteen tiivistelmäteksti kirjoitetaan opinnäytteen luettavan osan lomakkeen lisäksi myös pdf-tiedoston metadataan. Kirjoita tähän metadataan kirjoitettavaa teksti. Metadatatekstissa ei saa olla erikoismerkkejä, rivinvaiho- tai kappaleenjako-merkkiä, joten näitä merkkejä ei saa käyttää tässä. Jos tiivistelmäsi ei sisällä erikoismerkkejä eikä kaipaa kappaleenjako-merkkiä, voit hyödyntää makroa `abstracttext` luodessasi lomakkeen tiivistelmää (katso kommentti tässä TeX-tiedostossa alla). Metadatatiivistelmätekstin on muuten oltava sama kuin lomakkeessa oleva teksti.

Avainsanat Vastus, resistanssi, lämpötila

Preface

I want to thank Professor Pirjo Professor and my instructors Dr Alan Advisor and Ms Elsa Expert for their guidance.

I also want to thank my partner for keeping me sane and alive.

Otaniemi, 9 February 2023

Aarni O. Halinen

Contents

Abstract	3
Abstract (in Finnish)	4
Preface	5
Contents	6
Symbols and abbreviations	8
1 Introduction	9
1.1 Problem Statement	9
1.2 Thesis outline	10
2 Background	11
2.1 Zero trust architecture	11
2.2 Containerization and Docker	11
2.2.1 Linux containers	12
2.2.2 Docker	13
2.3 Kubernetes	14
2.3.1 Kubernetes objects	14
2.3.2 Kubernetes components	16
2.3.3 Admission controllers	17
2.3.4 Sidecar pattern in Kubernetes	18
2.4 Kubernetes network model	19
2.4.1 Container Network Interface	19
2.4.2 Calico	21
2.4.3 Cilium	22
2.4.4 Multus	23
2.4.5 Extended Berkeley Packet Filter	24
2.5 Example attack scenarios	26
3 Solution requirements	27
3.1 Threat modeling	27
3.2 Security requirements	29
3.3 Environment	29
4 Hardening Pod security	30
4.1 Restricted Pod Security Standard	30
4.2 Enforcing other best practices	32
4.2.1 Manual service account mounting	32
4.2.2 Enforcing the fields	34
4.3 Solution	35

5	Network Isolation	36
5.1	Applying firewall rules inside Pod network namespace	36
5.1.1	IPTables	37
5.1.2	Deployment	40
5.2	Own network namespace for sidecar	40
5.2.1	Creating isolated network for sidecar traffic	41
5.2.2	Mutating sidecars to multiple Pods?	44
6	Solution Evaluation	45
7	Discussion	46
8	Conclusion	47
	References	48
A	Dockerfile for penetration testing	52
B	Example webserver deployed with StatsD sidecar	53
C	Webserver and StatsD deployed with Multus networking	55

Symbols and abbreviations

Symbols

- ↑ electron spin direction up
- ↓ electron spin direction down

Operators

$\nabla \times \mathbf{A}$ curl of vector in \mathbf{A}

Abbreviations

API	Application Programming Interface
BGP	Border Gateway Protocol
BPF, cBPF	(classic) Berkeley Packet Filter
cgroups	Control groups
CLI	Command-line interface
CNCF	Cloud Native Computing Foundation
CNI	Container Network Interface
CRD	Custom Resource Definition
DAC	Discretionary Access Control
DoS	Denial of Service
eBPF	Extended Berkeley Packet Filter
gRPC	Google Remote Procedure Call
HTTP	Hypertext Transfer Protocol
IPC	Inter-process communication
K8s	Kubernetes
LXC	Linux Containers
MAC	Mandatory Access Control
NAT	Network Address Translation
NIC	Network Interface Controller
OOM	Out-of-memory
OpenVZ	Open Virtuozzo
OS	Operating System
PID	Process ID
RCE	Remote code execution
SELinux	Security-Enhanced Linux
Seccomp	Secure computing mode
TC	Traffic Control
VXLAN	Virtual Extensible LAN
XDP	eXpress Data Path
ZTA	Zero Trust Architecture

1 Introduction

During the last decade, the IT industry has shifted from monolithic software applications towards microservices. In the microservice architecture, each application is split into a suite of smaller independent services that handle part of the business logic [36]. Each service runs its own process, and the application data is sent between the components using lightweight communication mechanisms such as Hypertext Transfer Protocol (HTTP) and Google Remote Procedure Call (gRPC). This architectural approach also increases software agility because each service becomes an independent unit of development, deployment, operations, versioning, and scaling [49]. This modularity is often associated with benefits like faster delivery and improved scalability. **JB: The point is that microservices have become popular because of containers** During the same period of time, containers have become a widely used method to deploy applications. Containers are a lightweight virtualization technique in which the application is bundled with all its dependencies into a single deployable unit that executes on top of the host machine kernel [17]. As such, a single microservice is quite often built and deployed as a container. In more complex systems with multiple containers, an orchestrator is often used to manage workloads. Kubernetes is one of the most widely used orchestrator tools. **JB: Describe shortly what an orchestrator does and what is the difference between orchestrator and container runtime. Maybe this should be in the background.**

Similarly to how design patterns emerged from the birth of object-oriented software systems, the modularity of containers and microservices has allowed the development of distributed system design patterns [18]. The most common of these patterns is the sidecar pattern, in which peripheral tasks like logging and observability are split away as their own containers from the main application container. Basically, the sidecar pattern is an extension of the modularity of the microservices architecture inside the microservice itself. The benefits of sidecar pattern are similar to microservices, allowing better resource allocation, re-use of components in other services and provide a failure containment boundary, for example. In Kubernetes, the basic unit of deployment is called a Pod, which may include one or more co-scheduled containers. The Pod is analogous to a microservice, while it contains one main container and all its sidecars, bundled into one deployable unit. **JB: The sidecar pattern allows to add non-functional features to a container (e.g., logging, authentication, authorization) without them leaking into the source code (source code is only for business logic).**

1.1 Problem Statement

Although the sidecar pattern makes it easier to add peripheral tasks to applications, it opens up questions about application security. Quite often, developers rely on containers created by third parties for sidecar tasks. The source code of these sidecar containers is not always available and finding vulnerabilities therein is not a trivial task. Furthermore, malicious actors can use supply chain attacks and typo-squatting to trick victims into installing malicious sidecars to their clusters. Once malicious attackers gain access to the sidecar, any misconfigurations or permissive security

mechanisms put the whole cluster at risk. Furthermore, Kubernetes is not secure by default; on a fresh installation, most of the included security mechanisms are in permissive settings or outright disabled.

In Kubernetes, there is limited amount of security features available on container-level. Most of the security-related policies and capabilities are defined for the Pod, which essentially means that any capability required by the main application is inherited in the sidecar. Thus, any privileged workload, even in another container in the Pod, risks privilege escalation from the sidecar. In addition, Kubernetes' firewalling solution, Network Policies, are granted for the whole Pod instead of individual containers. Both of the aforementioned issues allow for lateral movement and further escalation for the attackers. Thus, any exploitable security issue in a sidecar container makes an optimal launchpad for attack against the whole cluster. **JB: Statements are too weak. Here, the reader should understand that sidecars may contain vulnerable (or potentially vulnerable) code, and there are NO mechanisms to protect containers running in the same Pod.**

Zero trust architecture and the principle of least privilege are common security paradigms for limiting lateral movement and further escalation in a system if any component within has been compromised. In both paradigms, the capabilities of an individual component are limited to only those that are required for the component to function. The capabilities, such as network access and any container privileges, are explicitly given to the component that requires them, while everything else is denied.

If these paradigms are successfully applied to sidecars, sidecars could only use operations and network access required, while anything else would be blocked. However, since Kubernetes provides limited security on container-level, we need to find some other ways to implement these paradigms inside the Pods. This thesis proposes a solution for restricting the capabilities of sidecar while minimally affecting the main container, thus improving the security by extending the paradigms within the Pod.

1.2 Thesis outline

The following chapter 2 gives background about containers, Kubernetes and explains their common attack vectors. It also discusses Kubernetes networking and container network interface plugins. Chapter 3 proposes ideas for isolating sidecars from the main application container. The chapter discusses both container and network security in the context of Kubernetes Pod. Chapter 6 introduces an implementation based on the findings of the previous chapter. The pros and cons of the solution are discussed in Chapter 7. Finally, chapter 8 discusses future research and concludes the thesis.

2 Background

2.1 Zero trust architecture

Conventional network security has historically focused on perimeter defense [50]. Subjects like workload resources and users inside the perimeter are often assumed to be trusted and implicitly given access inside the network, while any request originating outside the network is subject to more scrutiny. Although the systems seem initially secure, the modern IT landscape with cloud-based systems, third-party components, remote workers, etc. means that the perimeter can be breached via numerous attack vectors. **JB: Increase of the attack surface of threat actors.** Once any subject inside the perimeter is compromised, the attackers can gain access to all the resources that the subject is authorized to access and move laterally within the perimeter, escalating the attack on other resources.

Zero trust architecture (ZTA) is a security paradigm that focuses on data and resource protection and on the premise that trust must always be explicitly granted and continuously evaluated [50, 57]. In contrast to a single perimeter defense, the focus in ZTA is to create fine-grained access rules around each of the resources while at the same time enforcing rules that deny other access, which is not explicitly allowed. Following the principle of least privilege, the access rules are made as granular as possible so that the number of trusted subjects equals the actual number of subjects that require the access. This achieves a multi-layered security boundaries, where the breach of one component through the most outward perimeter does not compromise the whole system. Instead of having permission to access all resources within the perimeter, malicious actors could only laterally move to the resources that the compromised component required to function. Any other component is still protected by its own perimeter, which would require another successful attack to be breached. Thus, the compromised component is of limited usefulness to the attacker instead of serving as a general attack vector against the system. **JB: How is the Zero trust implemented in practice? Is there a relationship between zero trust and sidecars?**

2.2 Containerization and Docker

Figure 1 illustrates common virtualization models. While traditional virtualization techniques virtualize workloads on top of a hypervisor that shares hardware resources between virtual machines, containerization is a technique where virtualization occurs at the operating system level [53]. Processes executing in containers run on the kernel of the host machine. However, each container is isolated to its own network, process namespace, and so on; two containers on the same host OS do not know that they share resources. Furthermore, containers are similarly isolated from accessing host OS resources.

BSD jails and *chroot* can be considered early forms of containerization technology, so the idea of containers is not new [20]. Recent Linux container solutions rely on two main implementations: Linux Containers (LXC) -based solution that relies on kernel features such as control groups (cgroups) and namespaces, and a custom kernel



Figure 1: Virtualization models [20]

and Linux distribution called Open Virtuozzo (OpenVZ). Docker [33] is a hugely popular LXC-based container runtime and provides an easy-to-use API and tooling for creating and managing containers. Docker also provides containerization for other OSes as well. However, in this thesis we focus only on the Linux implementation.

2.2.1 Linux containers

The Linux containers technology implements container isolation and containment using a Linux kernel feature called namespaces [51]. Namespaces [29] are a construct that wraps a global system resource in an abstraction that makes it appear to the processes in the namespace that they have their own isolated instance of the global resource. There are a total of eight namespaces: i) *Cgroup* which is used for resource management, ii) *Inter-process communication* (IPC) that isolates POSIX message queues, etc., iii) *Network* which isolates network devices, stack ports, etc., iv) *Mount* for file system isolation, v) *Process ID* (PID), vi) *Time*, vii) *User* for isolating user and group identifiers, and viii) *UTS* which isolates hostnames and NIS domain names. For example, the *network* namespace provides each container with its own loopback device, and even *iptables* rules. In another example, *mount* namespace ensures that container has no visibility or access to the host's or other container's file system. Compared to other namespaces that concern the isolation of kernel data, *cgroups* focuses on limiting available system resources per namespace [51]. Each namespace can be configured with its own limits on CPU and memory usage and available devices. Using Docker as an example, setting `-cpu`, `-memory` and `-devices` options will limit available resources for the container.

Since all containers and the host machine run on the same kernel, any container that manages to breakout from isolation may compromise other containers, the host, and the whole kernel. To combat this container breakout, several security mechanisms

are adopted from the Linux kernel to restrict the capabilities of containers [51]. The mechanisms include Discretionary Access Control (DAC) mechanisms like Capability [28] and Secure computing mode (Seccomp) [30], and Mandatory Access Control (MAC) mechanisms such as Security-Enhanced Linux (SELinux) and AppArmor [1]. With Capability, the superuser (i.e. the root user) privilege is divided into distinct units, each of which represents a permission to process some specific kernel resources. The feature turns the binary "root/non-root" security mechanism into a fine-grained access control system, which makes it easier to follow the principle of least privilege. For example, processes like web servers that simply need to bind on a Internet domain privileged port (numbers below 1024) do not need to run as root; they can be granted with CAP_NET_BIND_SERVICE capability instead [34]. The Seccomp mechanism constrains which system calls a process can invoke. The available system calls are defined for a container through the Seccomp profile which is defined as a JSON file. The default Docker Seccomp profile [32] includes more than 300 system calls. SELinux is integrated into CentOS/RHEL/Fedora distributions and utilizes a label-based enforcement model, while AppArmor is available in Debian and Ubuntu distros and adopts a path-based enforcement model [51].

2.2.2 Docker

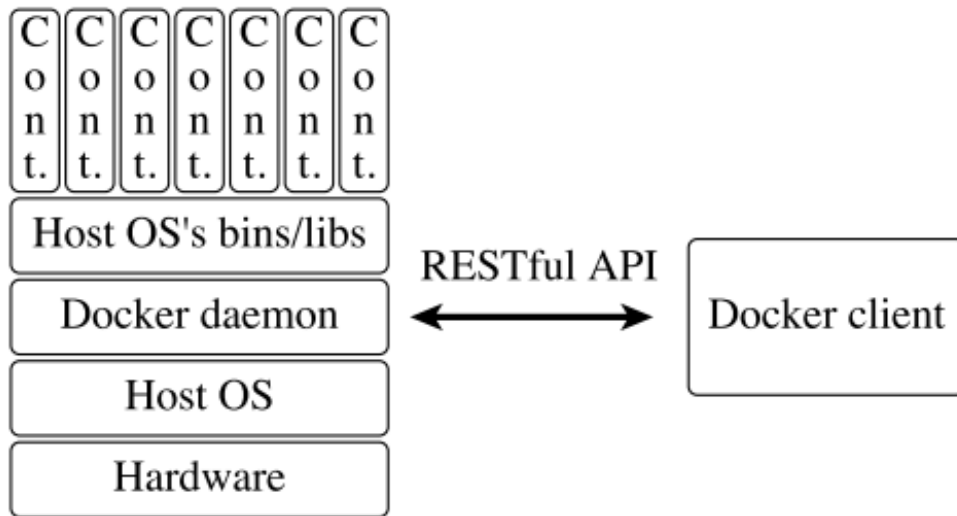


Figure 2: Architecture of Docker engine [17]

Docker is an open-source container technology written in Go and launched in 2013 [3, text]. The platform consists of the Docker Engine packaging tool, Docker image registries such as the Docker Hub public image repository and the Docker desktop application [2]. In general, the engine architecture is similar to container-based virtualization, as visible in figure 2 [17]. The containers run on top of the Docker

daemon, which manages and executes all the containers. The daemon is exposed to Docker clients via the RESTful HTTP API. The Docker client is a command-line tool which provides user interface for commanding the daemon, and thus containers. By exposing the API outside the host machine, the architecture enables remote control of the daemon with the client. Remote communication with the API should be carried out over TLS for security reasons.

Docker image is a read-only template with instructions for creating a Docker container [2]. Images are often based on another image, such as OS images `ubuntu` and `alpine`, with some additional customizations, such as the installation of web server binaries. Customizations are added to the image as a series of data layers so that each new command creates a new layer. This process makes image distribution more efficient since only the changes between layers must be distributed [17]. Layering is achieved with a special file system inspired by UnionFS which allows files and directories in different file systems to be combined into a single consistent file system.

Docker users can share their custom images publicly or privately in Docker Hub, or even host their own image registry platform. Most cloud providers also offer container registry services, so even proprietary software can be published in a private registry and used by other cloud services, like Kubernetes clusters. Whenever the image is not found locally, the client automatically tries to search and pull the image from the connected registries. **JB: Docker can only create and destroy containers. This makes it a runtime. This is an important concept to mention.**

2.3 Kubernetes

Kubernetes (K8s) [8] is an open-source container orchestrator, i.e., a system to automate the deployment, scaling and management of containerized applications. It allows the creation of a cluster which consists of a set of servers, called Nodes, on which application containers are scheduled by the system. The automation provides resilience and efficient resource utilization for workloads in the cluster: if a container or Node dies, the system tries to restart and reschedule containers so that the desired cluster state is maintained. K8s is hosted by the Cloud Native Computing Foundation (CNCF), but its origins are at Google where it was created as an open-source option for Google's proprietary Borg and Omega orchestrators [19]. K8s was open-sourced in 2014.

2.3.1 Kubernetes objects

Pods are the basic atomic scheduling unit in K8s. Pods consist of one or more tightly-coupled containers with shared storage volume and networking [12]. Containers in a pod are always co-located and co-scheduled and run in a shared context, i.e. a set of Linux namespaces. The network, UTS and IPC namespaces are shared by default, and the process namespace can be shared with `v1.PodSpec.shareProcessNamespace`. The common network namespace means that containers in a pod can communicate with each other via localhost, have a common IP address, and cannot reuse the same port numbers. In addition to the normal application container, Pods can include special

`initContainers` that are run only on Pod startup. These pods are used to modify the pod context before the actual workload starts. Multiple `initContainers` are run sequentially, and a failing container blocks the execution of the following initialization and normal workloads.

Instead of directly creating Pods, **Deployment** workload resources are used to create Pods in a cluster, even with singleton Pods [12]. With Deployments, the user describes the desired state in a declarative manner. The Kubernetes control loop then creates **ReplicaSet** based on the Deployment resource, which in turn guarantees the availability of the desired number of pods [7]. **DaemonSet** on the other hand is a workload resource that ensures that all or some Nodes run a copy of a Pod. Typical usecases for daemons are running Node monitoring and logging, and network plugins, which are discussed in depth in Chapter 2.4.1. **JB: If a Pod fails, there will not be a new one ready to be created. Sometimes developers want to have more control over the compute resources and specify a target state for them. This includes replication, scaling, and distribution across the various nodes. To account for these needs, K8s provides additional resources. Note that each of these resources will create one or more Pods in the end.**

All Pods across the cluster share the same subnet and can access each other via IP address. However, connecting to a Pod with IP address is suboptimal since Pods are ephemeral, and restarting a dead pod may receive a new IP address. Furthermore, horizontally scaled Pods with multiple replicas have as many IP addresses, making load balancing difficult.

Services are an object for exposing groups of Pods over a network [14]. The object defines a set of endpoints, that is, the targeted pods, along with a policy about how to make the pods accessible. The targeted pods are determined with a `selector` field in the object specification. Meanwhile, the `type` field determines how the Service is exposed. There are four different `ServiceTypes` levels **JB: Those are not levels but types.:** i) the default `ClusterIP` which exposes Service inside the cluster with its own IP address, ii) `NodePort` which exposes service in each Node's IP address on static port (by default within a range of 30000-32767), iii) `LoadBalancer` which exposes the Service externally using cloud provider's load balancer, and iv) `ExternalName` which is used to map Service to DNS name instead of a group of Pods. The field is designed as a nested functionality; each `ServiceType` level adds up to the previous one. Ingress object can also be used to expose services to outside the cluster. **JB: Usually services are proxies that route requests to the different pods. There are also Headless services which only add DNS entries for the pods.** The Ingress object requires installation of an Ingress Controller in the cluster. Cloud providers often have their own controllers, and all the examples in this thesis are executed on a local cluster where no external access is needed. **JB: Thus, the controllers are left as an exercise for the reader Do not discuss controllers here. You may add a paragraph on CustomResources and how one can easily extend K8s APIs. You could briefly discuss Ingresses and how they work.**

Namespaces provide isolation for cluster objects and allow grouping of objects under a single name. New K8s cluster starts with four namespaces: `default`, `kube-node-lease`, `kube-public` and `kube-system`. Namespaced objects like Deployments, Services and Pods are always deployed under a namespace which is

default if not explicitly defined. `kube-system` is the namespace for all objects created by the K8s system which is discussed in more detail in the next Chapter 2.3.2. Namespaces also provide a scope for naming; names of resources need to be unique within a namespace, but not across namespaces. Namespaces are also used to enforce resource quotas, access control, and isolation for cluster users, for example, in multi-tenancy setups. Pod Security Standards [11], which are used by the Pod Security admission controller, are also defined at the namespace level. Admission controllers are discussed in Section 2.3.3. **JB: Namespaces are only a logical grouping of resources. They do not provide isolation. They make it easy to provide isolation.**

2.3.2 Kubernetes components

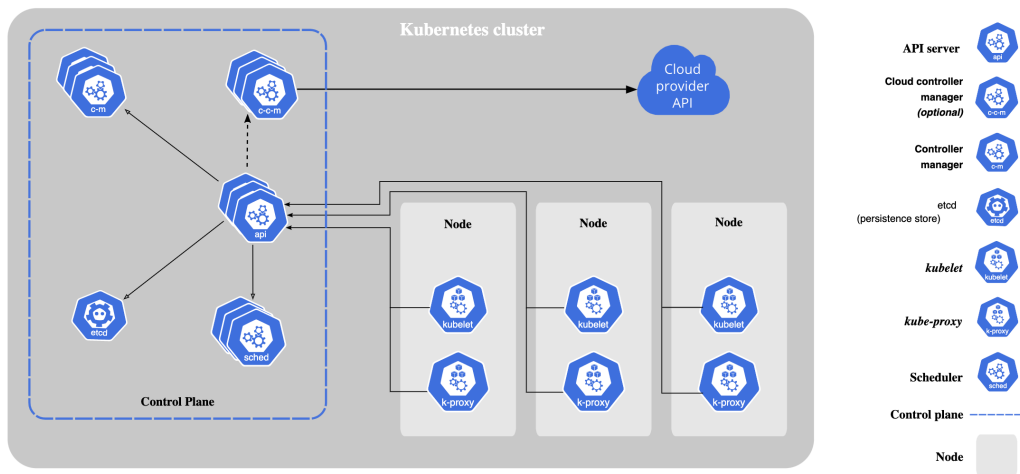


Figure 3: Kubernetes cluster architecture [9]

The figure 3 describes the Kubernetes cluster with a control plane and three worker nodes. The control plane consists of components that control, monitor, and store the state of the cluster; essentially, these are the components that are needed for the complete and working Kubernetes cluster [9]. The control plane components can run on any worker node. However, clusters often have a specialized master node for control plane components, which does not run any other containers. For fault tolerance and high availability, control plane components should run on multiple Nodes in production environments. The control plane consists of these main components:

API server is a front-end component of the control plane. It is a stateless **JB: HTTP is stateless by default but the server is not stateless. It holds the entire state of the application.** HTTP server that validates and authenticates commands given to the cluster. For valid commands, the server then forwards these to other control plane components. For example, the `kubectl` command-line interface (CLI) tool sends commands to the API server with HTTP. The main implementation of the server is `kube-apiserver`. The server can be horizontally scaled by running several instances on multiple Nodes and load-balancing traffic between the instances.

Etcd [4] is a highly consistent, distributed key-value store. It is the stateful component of the control plane: all of the cluster data is stored in etcd. Thus, the stability of the component is critical for the whole cluster. To tolerate failures, etcd implements a leader-based architecture. Multiple etcd clients automatically elect a leader instance as the source of truth. Other instances periodically update their state from the leader instance, so that the state stays eventually consistent across all instances. On leader failure, the other instances automatically elect a new leader to keep the system functioning.

Scheduler watches for newly created Pods that have no assigned worker node, and selects one of the active Nodes for them to run on [9]. Scheduling takes into account resource availability on Nodes, Pod resource requirements, object specification affinity rules and hardware, software, and policy constraints, among others.

Controller manager is a control plane component that runs all the controller loop processes [9]. Controller loops, like the Deployment controller, continuously watch the current and desired cluster state. When the states differ, they send commands via the API server so that the cluster moves towards the desired state. All the built-in controllers are compiled into a single binary, even though the controllers are logically different processes.

Each Node also has components that are essential for Kubernetes to work properly. **Kubelet** is an agent that ensures that containers are running in a Pod [9]. It receives a set of pod specification from the API server and ensures that containers are running on the Node, follow the pod specifications and are healthy. Kubelet only manages containers created by Kubernetes. **Kube-proxy** maintains network rules on Nodes. Part of the Service objects' networking is implemented by **kube-proxy**; the proxy writes iptables rules that route traffic [24].

2.3.3 Admission controllers

Admission controllers are a feature of the Kubernetes API server, used to validate and modify requests made to the server [5]. The controllers execute before the request is executed but after it is authenticated and authorized by the server. Several important features of Kubernetes are implemented with admission controllers, and these should be enabled on a properly configured API server. In addition to the built-in controllers, Kubernetes provides `MutatingAdmissionWebhook` and `ValidatingAdmissionWebhook` controllers for building own admission logic.

Admission controllers can be validating, mutating, or both [5]. Mutating controllers may modify related objects to the requests they admit, while validating controllers either approve or reject the request. The control process first executes the mutating controllers so that no mutations occur after the validation. If a controller in either phase rejects the request, the request is not processed further, and error is returned to the end-user.

One notable Admission controller is the `PodSecurity` controller. The controller validates the Pods before they are admitted, making sure that the requested Pod security context and other restrictions are permitted in the namespace to which the pod is assigned [5]. The controller is enabled by default, and can be taken into use just by

configuring Pod Security Admission labels for Namespace objects.

The labels use `pod-security.kubernetes.io/<MODE>: <LEVEL>` format, where `MODE` defines the action to be taken when the security level is violated and the `LEVEL` is a predefined level of the Pod Security Standard. The three available levels are `privileged`, `baseline` and `restricted` [11].

The available actions are i) `enforce`, which will reject the pod on violation, ii) `audit`, which triggers an event about the violation in the audit log, and iii) `warn`, which triggers user-facing warning about the violation [10]. A namespace can configure any or all three of the available modes and even set a different level for the modes. For example, it is possible to warn the user about violation of security policies without blocking the request by setting the `warn` mode more restrictive than `enforce`.

2.3.4 Sidecar pattern in Kubernetes

As mentioned before, Pods are the basic scheduling abstraction in Kubernetes and they support management and co-scheduling of multiple containers as an atomic unit. This co-scheduling and management of multiple symbiotic containers as a single unit enables multi-container application design patterns to emerge [18]. The sidecar pattern is the most common of these design patterns. As an example of this pattern, the main application container can be a simple web server paired with a container that collects server logs from a file and streams them to a centralized log management system. Another example of this pattern is the Istio service mesh [27] and its Envoy proxy sidecar, which routes all traffic through the Istio control plane for management, observability, and security operations.

In the pattern, peripheral tasks such as logging, configuration, and observability are isolated from the main application into helper containers. These containers, sidecars, are tightly-coupled to the parent application container and should share the lifecycle of the parent. Although sidecar functionality could be built into the main container, there are benefits in using separate containers [18]. The isolation allows tweaking of containers' cgroups so that CPU cycles can be prioritized for the main container. The isolation also provides a failure containment boundary between the main and sidecar processes. Since the container is also the unit of deployment, sidecar containers could be developed, tested, and deployed independently of each other. Sidecar containers can also be developed with different tools and dependencies and in a way that they are re-useable with other application containers. However, this multiplies the number of moving parts of the overall system, which increases the size of the test matrix considering all the combinations of container versions that might be seen in the production environment. **JB: The last sentence does not add new information. One thing to add here is how in k8s you setup sidecars. In most of the cases you either setup iptables rules to route traffic from the sidecar to the main container (and vice-versa) or use shared files to read information (e.g., logs).**

2.4 Kubernetes network model

A fundamental part of a Kubernetes cluster is how nodes and resources are networked together. Specifically, the networking model needs to address four different types of networking problems: i) intra-Pod (i.e., container-to-container within same Pod) communication, ii) inter-Pod communication between Pods, iii) Service-to-Pod communication and iv) communication from external sources to Services [6]. The model also requires that each Pod is IP addressable and can communicate with other Pods without network address translation (NAT), even when Pods are scheduled on different hosts [56]. All agents on a host should also be able to communicate with Pods on the same host. The implementation of this model is not part of Kubernetes, but is handed to special plugins that implement the Container Network Interface (CNI) specification. **JB: K8s defines the CNI interface to allow different vendors and operators to use different networking mechanisms in kubernetes.**

2.4.1 Container Network Interface

The Container Network Interface (CNI) [15] is a networking specification, which has become the de-facto industry standard for container networking. It is backed by CNCF [56]. CNI was first developed for the container runtime rkt [46], but it is supported by all container runtimes, and there are a large number of implementations to choose from. Most container orchestrators have adopted the specification as their networking solution. The biggest outlier is Docker Swarm, which instead implements libnetwork [35]. **JB: This may be not needed**

The CNI specification has five distinct definitions: i) a format for network configuration, ii) an execution protocol between the container runtimes and the plugin binary, iii) a procedure for the runtime to interpret the configuration and execute the plugins, iv) a procedure for delegating functionality between the plugins, and v) data types for the plugins to return their results to the runtime [15]. The network configuration is defined as a JSON file, and it includes a list of plugins and their configuration. The container runtime interprets the configuration file at the plugin execution time and transforms it into a form to be passed to the plugins. The execution protocol defines a set of operations (ADD, DEL, CHECK, VERSION) for adding and removing containers from the network. The operation command, similarly to other protocol parameters, is passed to the plugins via the OS environment variables. The configuration file is supplied to the plugin via stdin. On successful execution, the plugin returns the result via stdout with a return code of 0. On errors, the plugin returns a specific JSON structure error message to stderr and a non-zero return code. When runtime mutates a container network, it results in a series of ADD, DELETE, or CHECK executions. These are then executed in the same order as defined in the plugins list, or reversed order for DELETE executions. Each plugin then returns either Success or Error JSON object. The execution of a series of operations ends when it encounters the first error response, or when all the operations have been performed.

The CNI plugin must provide at least connectivity and reachability for the containers

[45]. For connectivity, each Pod must have a network interface controller (NIC) for communication outside its networking namespace. The NIC must have an IP address that is reachable from the host Node, so that cluster processes like Kubelet health and readiness checks can reach the Pod. Reachability means that all Pods can be reached from other Nodes directly without NAT. Thus, each pod receives a unique IP address from the PodCIDR **JB: This should be introduced properly** range configured on the Node by the Kubelet bootstrapping phase. The end-to-end reachability between different Node PodCIDRs is established by encapsulating in the overlay network, for example, with Virtual Extensible LAN (VXLAN), or orchestrating on the underlay network, e.g. with Border Gateway Protocol (BGP).

Since Kubernetes does not provide networking between Pods, it has no capabilities to enforce network isolation between workloads. Thus, another key feature of CNI plugins is to enforce network traffic rules. Kubernetes provides a common object called `NetworkPolicy` for CNI plugins to consume. The listing 1 is an example of network policies for a web application front-end. The first policy functions as a default rule that denies all implicit egress and ingress traffic. The second policy allows traffic to the default HTTP port, and the third policy allows the front-end application to send traffic to any Pod in the *back-end* namespace.

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: deny-all
5   namespace: frontend
6 spec:
7   podSelector: {}
8   policyTypes:
9     - Ingress
10    - Egress
11 ---
12 apiVersion: networking.k8s.io/v1
13 kind: NetworkPolicy
14 metadata:
15   name: webapp-ingress
16   namespace: frontend
17 spec:
18   podSelector:
19     matchLabels:
20       app: webapp
21   policyTypes:
22     - Ingress
23   ingress:
24     - ports:
25       - protocol: TCP
26         port: 80
```

```

27 ---
28 apiVersion: networking.k8s.io/v1
29 kind: NetworkPolicy
30 metadata:
31   name: webapp-egress
32   namespace: frontend
33 spec:
34   podSelector:
35     matchLabels:
36       app: webapp
37   policyTypes:
38     - Egress
39   egress:
40     - to:
41       - namespaceSelector:
42         matchExpressions:
43           - key: namespace
44             operator: In
45             values: ["backend"]

```

Listing 1: Example NetworkPolicies for a frontend web application

The NetworkPolicy specification consists of a `podSelector` that specifies pods that are subject to the policy and `policyTypes` to specify the Ingress and Egress rules for the traffic [16] to the target Pod. Each rule includes `to` or `from` field for selecting Pod, Namespace or IP address block in CIDR notation on the other side of the connection, and `ports` field for explicitly specifying which ports and protocols are part of the rule. The policies are additive; when multiple rules are defined for a Pod, the traffic is restricted to what is allowed by the union of the policies. Many CNI plugins also introduce Custom Resource Definitions **JB: if you add CRD here then you should discuss them in the earlier section** for their own, more granular, network policy rules. **JB: Not all CNIs implement NetworkPolicy rules (e.g., flannel)**

Although all CNI plugins meet the requirements listed above, they may differ in architecture significantly. The plugins can be classified based on which OSI model network layers they operate on, which Linux kernel features they use for packet filtering and which encapsulation and routing model they support for inter-host and intra-host communication between Pods. In this thesis, we focus on three different CNI plugins: Calico, Cilium, and Multus.

2.4.2 Calico

Calico [59] is an open-source CNI plugin with modular architecture that supports a wide range of deployment options. Each Pod created in the Calico network receives one end of a virtual ethernet device link as its default `eth0` network interface, while the other end is left dangling on the host Node [43]. The Pod end of the link receives an IP address from Pod CIDR, but the Node end does not. Instead, a `proxy_arp`

flag is set on the host side of the interface while containers have a route to link-local address 169.254.1.1, thus making the host behave like a gateway router. For routing packets between Nodes, Calico creates a VXLAN overlay network. Optionally, Calico supports IP-in-IP overlay or non-overlay network with BGP protocol.

On each Node, a `calico-node` daemon setups CNI plugin, IPAM, and possible eBPF programs. The daemon subscribes to Kubernetes API for Pod events and manages both container and host networking namespaces. Calico also deploys a single-container `calico-kube-controllers` Pod into the Kubernetes control plane. The container executes a binary that consists of controller loops for Namespace, NetworkPolicy, Node, Pod, and ServiceAccount Kubernetes objects. The Calico project also introduces its own CLI tool, called `calicoctl` [61], to manage Calico's custom resources. The tool provides extra validation for the resources which is not possible with `kubectl`.

Calico supports Kubernetes NetworkPolicies as well as its own namespaced `projectcalico.org/v3.NetworkPolicy` Custom Resource Definition (CRD). Both of the policies work on OSI layers L3 (identity, e.g. IP address) and L4 (ports). Compared to the built-in policy, the Calico policy includes features such as policy ordering, log action in rules, and more flexible matching criteria (e.g., matching on ServiceAccounts) [60]. The policy can also match other Calico CRDs such as **HostEndpoints** and **NetworkSets**, which allows implementing rules on host interfaces and non-Kubernetes resources. If Calico is installed along the Istio service mesh, the Calico Network Policy can enforce L7 (e.g. HTTP methods and URL paths) policies on the Envoy proxy. For policies that are not tied to a Kubernetes namespace, Calico provides a `GlobalNetworkPolicy` CRD.

JB: Now that Calico uses eBPF there is little difference between them and both can reduce the amount of sidecars to one per node.

2.4.3 Cilium

Cilium [21] is one of the most advanced and powerful CNI plugins for Kubernetes. Similarly to Calico, it creates a virtual ethernet device for each Pod and sets one side of the link into the Pod's network namespace [44] as the default interface. Cilium then attaches extended Berkeley Packet Filter (eBPF) programs to ingress traffic control (tc) hooks of these virtual ethernet devices to intercept all incoming packets from the Pod. The packets are intercepted and processed before the network stack, and thus `iptables`, reducing latency 20%-30% and even doubling the throughput of packets in some scenarios [16]. The network between Pods running on different hosts is handled by default with VXLAN overlay, but there is support for Geneve interfaces and native-routing with the BGP protocol as well [21].

The Cilium system consists of an agent (`cilium-agent`) daemon running on each Node, one or more operator (`cilium-operator`) Pods and a CLI client (`cilium`) [22]. The agent daemons subscribe to events from the Kubernetes API and manage containers' networking and eBPF programs. The CLI tool, which is installed on each agent, interacts with the REST API of the agent and allows one to inspect the state and status of the local agent. The tool should not be confused with the Cilium management

CLI tool, also incidentally named `cilium`, which is typically installed remote from the cluster. The operator is responsible for all management operations that should be handled once for the entire cluster, rather than once for each Node. This includes, for example, the registration of CRDs.

Although default Kubernetes NetworkPolicies provides security on OSI layers L3 and L4, Cilium provides CRDs that also support L7 policies [23]. If L7 policies exist, the traffic is directed to Envoy instance bundled into the agent Pod which filters the traffic. Unlike on layers 3 and 4, policy violation does not result in dropped packet but an application protocol specific denied message. For example, HTTP traffic is denied with HTTP 403 Forbidden and DNS requests with DNS REFUSED. Cilium provides `CiliumNetworkPolicy` CRD that supports all L3, L4, and L7 policies. Cilium also provides `CiliumClusterwideNetworkPolicy` custom resource which is used to apply network rules to all namespaces in the cluster or even to nodes when using `nodeSelector`.

As even more advanced features, Cilium also includes natively `kube-proxy` replacement, encryption for Cilium-managed traffic, and Service Mesh, among others. By default, `kube-proxy` uses `iptables` to route the Service traffic [24]. With `kubeProxyReplacement` installation option, Cilium implements Service load-balancing as XDP and TC programs on the Node network stack. For encryption, Cilium supports IPsec and WireGuard implementations [25]. The service mesh performs a variety of features directly in eBPF, thus functioning without sidecar containers or proxying requests through the agent Pod's Envoy [37]. Since all features are not available as eBPF programs or on all kernel versions, Cilium automatically probes the underlying kernel and automatically reverts to the Envoy proxy when needed. For capabilities beyond the built-in mesh, Cilium also provides integration with Istio.

JB: Have you thought about how your solution is impacted by eBPF?

2.4.4 Multus

Traditionally, CNI plugins provide only a single network interface for a pod, apart from the loopback device. Multus [41] is a CNI plugin that allows the attachment of multiple network interfaces for a Pod. It does not provide any connectivity or reachability for containers like other plugins. Instead, it is installed as the first plugin in the CNI plugin chain. When executed, the plugin delegates the creation of the interface to other installed plugins. Since Multus does not provide any networking and thus does not independently, it is often called *meta plugin* to distinguish it from common CNI plugins like the previous Calico and Cilium.

Multus system includes a binary, a CNI configuration file, and a namespaced `NetworkAttachmentDefinition` CRD that is used to define network interfaces used in Pods. The binary and the configuration file are often installed on cluster nodes via a `DaemonSet`. The daemon consists of an *initContainer* that copies the binary into the `/opt/cni/bin` directory, and a daemon container that configures the configuration file and optionally spawns an HTTP server for additional features such as metrics [41]. The configuration file satisfies the CNI specification with few additional attributes of which the combination of `clusterNetwork` and

`defaultNetworks` or `delegates` are imperative for the CNI plugin to function [42]. The `clusterNetwork` specifies the main network of the cluster, which implements the `eth0` interface and the Pod IP address. The `defaultNetworks` is an optional array of networks that should be added for any Pods by default. The values can be names of the `NetworkAttachmentDefinition` objects or paths to the CNI plugin's JSON configuration files. Optionally, the `delegates` attribute can be used; it supports similar format of values. In this scenario, the first element of the array functions as `clusterNetwork` and the rest are inferred as `defaultNetworks`.

Attaching additional interfaces to workloads is most often configured by adding a special annotations field `k8s.v1.cni.cncf.io/networks` to workload resource definitions. In the simplest configuration, the field takes a comma separated list of the `NetworkAttachmentDefinition` names as input. The network interface identifiers can be modified by providing the attachment input in `name@interface-identifier` format. Otherwise, Multus names the interfaces `net0`, `net1` and so on. If extra configuration for the networks is needed, the annotation also supports the JSON array format.

2.4.5 Extended Berkeley Packet Filter

Berkeley Packet Filter (BPF, or nowadays often cBPF) was originally developed in the early 1990s as a high-performance tool for user-space packet capture [52]. BPF works by deploying the filtering part of the application, `packet filter`, in the kernel-space as an agent. The `packet filter` is provided with a program (often denoted as BPF program) consisting of BPF instructions, which works as a set of rules for selecting which packets are of interest in the user-space application and should be copied from kernel-space to user-space. The instructions are executed in a register-based pseudo machine. Since network monitors are often interested only in subset of network traffic, this limits the number of expensive copy operations across the kernel/user-space protection boundary only to packets that are of interest in the user-space application. A notable use-case for BPF is `libpcap` library, which is used by a network monitoring tool called `tcpdump`.

Later in the 2010s the Linux community realized that BPF and its ability to instrument the kernel could benefit other areas than packet filtering as well [62]. This reworked version of BPF was first merged into the Linux kernel in 2014 and is publicly called the extended Berkeley Packet Filter (eBPF) to distinguish it from the original cBPF. The kernel development community continues to call the newer version BPF, but instead of the original acronym consider it a name of a technology. Similarly to the kernel community, the term BPF always refers to the eBPF in this thesis.

The eBPF programs are compiled to bytecode and loaded into the kernel with `bpf()` system call [54]. Most often, programs are written in restricted C and compiled with the LLVM Clang compiler to bytecode. It is also possible to use the eBPF assembly instructions and `bpf_asm` utility for converting instructions to bytecode. eBPF programs follow an event-driven architecture: a loaded eBPF program is hooked to a particular type of event, and each occurrence of the event triggers the program execution.

There are two different network event interfaces in eBPF: `eXpress Data Path`

(XDP) and Traffic Control (TC) [54]. XDP programs are attached to an NIC and can handle only incoming packets [47]. The programs are called directly by the NIC driver if it has XDP support, thus executing before packets enter the network stack. This skips expensive packet parsing and memory allocation operations and allows XDP programs to run at very high throughput. Thus, even the main network buffer *skbuff* is not populated. Some SmartNICs even support offloading the program to the NIC's own processor from the host CPU, further improving the performance of the host machine [26]. If the driver does not support XDP, generic XDP is used and the programs run after the packet has been parsed by the network stack.

XDP programs can read and modify the contents of packets [62]. Since the packets are not parsed into the network stack, the programs have to work with raw packets and implement their own parsing functionality. The program's return value determines how the packet should be processed further. With `XDP_DROP` and `XDP_PASS` return values, the packet can be dropped or passed further to the networking stack, respectively. The packet can also be bounced back to the same NIC on which it arrived with `XDP_TX`, usually after modifying the contents of the packet. `XDP_REDIRECT` is used for redirecting the packet to a different NIC, CPU or even to another socket.

TC programs are executed when both incoming and outgoing packets reach the kernel traffic control function within the Linux network stack [62]. The ingress hook runs after the packet is parsed to *skbuff* but before most of the network stack. On egress the stack is traversed in reverse; thus, the hook executes after most of the network stack. TC programs can read and write directly to a packet in memory. Similarly to XDP programs, the return value of the program determines the further processing of the packet. The packet can be passed further in the stack with `TC_ACT_OK`, dropped with `TC_ACT_SHOT`, or the modified packet can be redirected back to the start of the classification with `TC_ACT_RECLASSIFY`, among others. JB: Do we need this level of detail to understand your solution?

2.5 Example attack scenarios

JB: This is part of the threat modeling

- Privileged container
- CAP_SYS_ADMIN, mounting /proc and chroot
- CAP_SYS_PTRACE, shellcode injection to running program, nc 172.17.0.1 on port running shell
- Mounted docker socket, creating privileged containers

All of the example attack scenarios start by an attacker getting shell access to a container running in a Pod, usually through a remote code execution (RCE) flaw on the container application and then executing reverse shell inside the Pod. The scope of the examples is not in the initial attack, but in the Pod template misconfigurations that then provide some path for the attacker to escalate the attack and even take over the whole cluster in the end.

TODO: add file

The first attack scenario includes a privileged container, as described in File 2.5. Privileged containers have all the capabilities of the host machine, so practically they can be perform almost any action available on the host. This includes, but is not restricted to, for example "pipeing" (TODO: a better word for undock.sh and other scripts) commands as root to the Node's shell and mounting the host's disks. If combined with `hostPID: true`, the attacker can see all processes on the host and use `nsenter` to execute commands in the namespaces of the other processes.

TODO: add file

The second example file 2.5 does not have similar privileges for executing commands as the first, but has unlimited access to mount the entire host's filesystem, with read and write access. Thus, the attacker can try to find any credentials stored on the host machine and use these to escalate the attack. Important credentials include `kubeconfig` files that store access tokens to the K8s API server, `ServiceAccount` tokens that may have been mounted on any Pod on the host, SSH keys and hashed user passwords in `/etc/shadows`.

TODO: some extra issues (finding `.kubeconfig` files, running on control-plane -> `etcd` and secrets within, cloud metadata)

3 Solution requirements

JB: Threat model can also be a separate section. In principle, you should discuss the attacks you want to defend against (e.g., installing a malicious sidecar, or legitimate sidecar with vulnerabilities), the location of the attacker in the system (inside a specific container, in any container, outside the cluster), the prerequisites for an attack, and which parts of the threat matrix are involved.

This chapter analyzes threats in the K8s cluster from the perspective of sidecars and defines the security requirements for the solution introduced in the following chapters. First, threat modeling is used to identify possible attack vectors and readily available mitigations. Then, the model is used to identify all issues that should be solved with the solution. Finally, an example development environment is introduced to test out the solution.

3.1 Threat modeling

A security threat is any possible event in a system that could lead to a potential loss of confidentiality and integrity of an asset in the system. Threat modeling is a structured approach to identify and prioritize potential threats to a system. This includes profiles of potential attackers and their goals and methods, as well as potential mitigations [58]. The idea is to build a catalog of all potential threats, prioritize issues and try to find mitigations so that attackers cannot use the vulnerabilities against the system.

For generic K8s clusters there exist a number of extensive threat models, for example the *Threat Matrix for Kubernetes* by Microsoft [55]. The scope of this thesis is on sidecar security, so we will only focus on the possible threats that can be leveraged after the attacker has gained access to a sidecar container inside the cluster. The *Threat Matrix for Kubernetes* is used to identify threats in this scenario, while the results are listed in Table 1.

Table 1: K8s sidecar threat model

Threat	Description	Mitigation
Privileged containers	If containers are given privileges, malicious actor can breakout from the container and escalate the attack on cluster.	Use restricted Pod Security Admission to enforce security rules. PSAs are defined for namespaces, so isolate privileged containers into their own namespaces to follow the principle of least privilege.
Writing to the host file system	TODO.	Do not allow unnecessary mounts. Use <code>readOnlyRootFilesystem: true</code> whenever possible.
Permissive RBAC on service accounts	Containers in a Pod share service accounts ¹ . By default, Pods automatically mount the service account to all containers.	Disable automatic mounting of service accounts ² and explicitly mount them to containers when needed. Adhere to the principle of least privilege.
There are no resource limits for containers	If not set, a single container can hog system resources, causing denial of service (DoS).	Set resource limits to containers ³ .
The Pod has network access to the control plane	By default, all traffic inside the cluster is allowed.	Use NetworkPolicies as a firewall. Deny by default and explicitly allow traffic only if needed.
Sidecar can use the main container's NetworkPolicy to access kube-system or other Pods	NetworkPolicies affect the NIC of the Pod network, which is shared by sidecars.	No built-in solution available!
Sidecar has network access to cloud resources	Cloud providers may attach cloud identities to cluster Nodes, so that Pods can access resources outside the cluster.	Same issue as above! Block mounting volumes with access to cloud credentials.
Sidecar has network access to the main container	Main container is accessible via loopback device, which cannot be protected with NetworkPolicy.	There is no built-in solution available!
Sidecar is used to sniff the main application's traffic	TODO	mTLS? Requires a service mesh.
Unencrypted secrets	TODO	Encrypt at rest

3.2 Security requirements

Threat modeling identified two main categories of issues: permissive workload configurations and networking related issues. Essentially, all the threats are caused by sidecars not respecting the principle of least privilege; the sidecar inherits execution and networking privileges from the main application container. Based on the model, the solution should provide answers to these main questions:

1. How to ensure and enforce that no workloads that conflict with the principle of least privilege are deployed to the cluster?
2. How to enforce Zero Trust network that allows traffic filtering on container-level and limits communication on both loopback and Pod network interface device?

As for the first question, existing mitigations were already found while threat modeling. A solution in which all mitigations are applied is introduced in Chapter 4. For the second question on building the Zero Trust network, the Pod must be firewalled for both inter-Pod communications on the Pod network NIC and intra-Pod communications on the loopback NIC. However, K8s does not provide any built-in solution for creating these firewalls. Few possible solutions for building the Zero Trust network inside the Pod are given in the Chapter 5.

3.3 Environment

The solution is tested and developed on a Minikube cluster running on a local machine and using Docker as the driver for Nodes. The cluster consists of two Nodes, the purpose of which is to deploy control plane components separately from worker ones. Furthermore, the setup also allows testing of network between components hosted on different Nodes. The complete setup is hosted on Github (TODO: add link). Both Calico and Cilium are used as CNI plugins, since the selection of CNI plugin has minor implications on the actual networking solution. The implications are discussed in detail in chapter 5.

A simple Node.js webserver with StatsD sidecar is used as the application workload. The source code can be found in Appendix B. The webserver serves "Hello world" on port 8888 and writes response times to the StatsD sidecar via the loopback device on port 8125, which is the default for StatsD.

For penetration testing purposes, another sidecar is introduced with common networking and K8s command line tools. When deployed instead of the StatsD client, this sidecar can be used to simulate situations where the attacker has managed to get access to the shell inside the sidecar. The Dockerfile for the image can be found in Appendix A.

4 Hardening Pod security

This chapter provides a solution for hardening Pods against privilege escalation attacks. The solution enforces that deployed resources follow K8s best security practices. Most of the practices are enforced by the built-in Pod Security Admission controller. However, this chapter also introduces extra security measures that fix a few oversights regarding sidecar containers in the controller.

4.1 Restricted Pod Security Standard

Since version 1.25, K8s has shipped with the Pod Security Admission controller as a stable feature. The controller, as discussed in Chapter 2.3.3, provides three different Pod Security Standards that can be used to warn and enforce against insecure Pod configurations. The restricted Pod Security Standard is the strictest of the standards and aims at the best Pod hardening practices [11], so it will be the most optimal for our solution. The table 2 lists all fields affected by the standard.

Table 2: Pod fields enforced by restricted Security Standard

Field name	Usage	Allowed values
hostPID, hostIPC, hostNetwork	Controls whether container uses host's PID, IPC and network namespace.	false
privileged	Controls whether Pod can run privileged containers.	false
capabilities.add	Defines Linux capabilities for the container.	NET_BIND_SERVICE
capabilities.drop	Defines Linux capabilities for the container.	ALL
volumes[*]	All volume types are not allowed. For example, hostPath, that maps host directories, are not allowed.	volumes[*].configMap, volumes[*].csi, volumes[*].downwardAPI, volumes[*].emptyDir, volumes[*].ephemeral, volumes[*].persistentVolumeClaim, volumes[*].projected, volumes[*].secret
hostPort	Expose container via host's network port.	undefined
container.apparmor.security.beta.kubernetes.io/* annotation	Sets the AppArmor profile used by containers. On supported hosts, the runtime/default AppArmor profile is applied by default.	runtime/default, localhost/*
seLinuxOptions	Sets the SELinux context of the container.	Set if supported by environment.
procMount	The default /proc masks are set up to reduce attack surface, and should be required.	Default
seccompProfile.type	Sets the seccomp profile used to sandbox containers.	RuntimeDefault or Localhost
sysctls[*].name	Sysctls can disable security mechanisms or affect all containers on a host, and should be disallowed except for an allowed "safe" subset.	kernel.shm_rmid_forced, net.ipv4.ip_local_port_range, net.ipv4.ip_unprivileged_port_start, net.ipv4.tcp_syncookies, net.ipv4.ping_group_range
allowPrivilegeEscalation	Restricts escalation to root privileges.	false
runAsNonRoot	Controls whether container can run as root user.	true
runAsUser, runAsGroup	Controls the user and	Set both to non-zero

```

1 apiVersion: v1
2 kind: Namespace
3 metadata:
4   name: foo
5   labels:
6     name: foo
7     pod-security.kubernetes.io/enforce: restricted
8     pod-security.kubernetes.io/audit: restricted
9     pod-security.kubernetes.io/warn: restricted
10    pod-security.kubernetes.io/enforce-version: latest
11    pod-security.kubernetes.io/audit-version: latest
12    pod-security.kubernetes.io/warn-version: latest

```

Listing 2: Namespace resource with restricted Security Standard

The Security Standard can be enforced trivially by adding `pod-security.kubernetes.io/enforce: restricted` as a label on a K8s namespace resource, as shown in Example ??.

4.2 Enforcing other best practices

The restricted Security Standard hardens the Pod against most of the identified security threats. However, it still does not enforce specific resource limits and allows automatic mounting of Service Accounts for the Pod containers.

Adding resource limits to containers is straightforward: just add values to both `resources.limits.cpu` and `resources.limits.memory` fields for all of the containers, similarly to Example 4. The CPU usage is measured in CPU units and can also be expressed in millicpus, ie. both "1000m" and integer value of 1 are equivalent to 1 physical or virtual core [13]. For memory, the base unit is bytes, but it also supports quantity suffixes like "M", "Mi", and "Gi" for megabytes, mebibytes, and gigibytes, respectively. The resource limits are registered to container's cgroup by the Kubelet. The limits are hard, which means that if a container exceeds its CPU limit, the execution is blocked until more CPU capacity is available. Exceeding the memory limit causes termination with an out-of-memory (OOM) error.

4.2.1 Manual service account mounting

By default, the Service Account tokens are mounted in the `/var/run/secrets/kubernetes.io/serviceaccount` directory in every container. This feature can be disabled by setting `automountServiceAccountToken: false`, but then any container that actually uses the Service Account must receive the token in some other way. Since volume mounts are defined per container and service account tokens can be created manually with Secrets, the issue can be circumvented by manually mounting the token to containers that use it.

```

1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: Role
3 metadata:
4   namespace: foo

```



```

5   name: foo-service-account-role
6   rules:
7   - apiGroups: [""] # "" indicates the core API group
8     resources: ["pods"]
9     verbs: ["get", "watch", "list"]
10  - apiGroups: [""] # "" indicates the core API group
11    resources: ["pods"]
12    verbs: ["get", "watch", "list"]
13  ---
14  apiVersion: v1
15  kind: ServiceAccount
16  metadata:
17    name: foo-service-account
18    namespace: foo
19  ---
20  kind: RoleBinding
21  apiVersion: rbac.authorization.k8s.io/v1
22  metadata:
23    name: foo-service-account-rolebinding
24    namespace: secure-ns
25  subjects:
26  - kind: ServiceAccount
27    name: foo-service-account
28    apiGroup: ""
29  roleRef:
30    kind: Role
31    name: foo-service-account-role
32    apiGroup: ""
33  ---
34  apiVersion: v1
35  kind: Secret
36  metadata:
37    name: foo-service-account-token
38    namespace: foo
39    annotations:
40      kubernetes.io/service-account.name: foo-service-
        account
41  type: kubernetes.io/service-account-token

```

Listing 3: ServiceAccount with permissions for fetching Pod resources

Example 3 shows how to create a Service Account token with permission to read the status of Pods in the cluster. While Role, RoleBinding, and ServiceAccount resources are used to create and define RBAC rules for the Service Account, the Secret resource creates a similar authorization token as Pod with auto-mounting of service account tokens would create and mount on the containers. When the token created

by the Secret resource is mounted to the same path as the automatic mounting, as in Example 4, the token is only mounted to the container that needs it. With this approach, the account tokens are only mounted into containers that actually need them, and the tokens cannot be used for privilege escalation from other containers of the Pod.

```
1 apiVersion: apps/v1
2 kind: Deployment
3 spec:
4   replicas: 1
5   selector:
6     matchLabels:
7       app: app
8   template:
9     spec:
10      automountServiceAccountToken: false
11      containers:
12        - name: main
13          image: main-image
14          resources:
15            limits:
16              cpu: "2"
17              memory: 1024Mi
18          volumeMounts:
19            - name: service-account
20              mountPath: /var/run/secrets/kubernetes.io/
21                serviceaccount
22        - name: sidecar-container
23          image: sidecar-image
24          resources:
25            limits:
26              cpu: "1"
27              memory: 512Mi
28      volumes:
29        - name: service-account
30          secret:
31            secretName: foo-service-account-token
```

Listing 4: Two container Pod with resource limits and manually mounted ServiceAccount

4.2.2 Enforcing the fields

There is no built-in enforcement tool for the fields in Kubernetes. However, the Admission Controller can be used for enforcement by creating a custom ValidatingAdmissionWebhook that fails Pod creation if the fields are not correctly set. This is

actually how the Pod Security Admission works under the hood. It is just built-in to the Kubernetes system as an original hook.

4.3 Solution

TODO: Summarize all the solutions into one.

5 Network Isolation

Implementing a zero trust network architecture in K8s cluster is not trivial, since building network isolation between containers in the same Pod is not possible with common CNI plugins and Network Policies. The root cause for this is that the lowest level of networking abstraction in K8s is the Pod, and that by definition, all containers in a Pod share network namespace. Thus, all traffic outside the Pod, from any of the containers, passes through the same `eth0` NIC and shares a common source IP address. Since NetworkPolicies operate on L3/L4, they cannot distinguish whether traffic originates from the sidecar or from the main container. Although ingress traffic can be easily identified with the destination port number, the source port for egress traffic depends on the TCP implementation. As a result, NetworkPolicies cannot be used to manage egress traffic from sidecar independently of the main container. For intra-Pod communication, the traffic goes through the loopback device which is not managed by CNI plugins. This means that NetworkPolicies do not have an effect on the loopback device. Furthermore, since all intra-Pod traffic has source and destination IP addresses of `127.0.0.1`, the packets have no clear identifiers that could be used for traffic filtering. This chapter introduces two general approaches as solutions for the issue.

The first solution enforces firewall rules inside the application Pod's network namespace, while the second approach creates network namespaces for each container by deploying them to their own Pods. While deploying containers in their own pods is an obvious option for creating the isolation, the approach breaks the tight integration of sidecar containers to the application container. Most notably, the sidecars will have their own independent lifecycle and scheduling, and the communication via the loopback device will be broken. In the scope of this thesis, we consider sidecars to be scheduled along the application container and be accessible via `127.0.0.1`. Thus, re-introducing these characteristics to the containers is also part of the solution.

5.1 Applying firewall rules inside Pod network namespace

Figure 4 illustrates the network architecture with all required network rule enforcement points. The solution needs to enforce network rules in four scenarios: i) external components and Pod's `eth0`, ii) main container and `lo`, iii) sidecar container and `lo`, and sidecar container and `eth0`. The ingress and egress rules in the first scenario are more easily enforced with NetworkPolicies, as shown in Listing 5. The policies allow cluster and external traffic to the main application port 8888, while denying all other traffic into and out of the Pod.

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: deny-all
5   namespace: app
6 spec:
```




Figure 4: Single network namespace architecture

```

7  podSelector: {}
8  policyTypes:
9    - Ingress
10   - Egress
11 ---
12 apiVersion: networking.k8s.io/v1
13 kind: NetworkPolicy
14 metadata:
15   name: web-allow-external
16   namespace: app
17 spec:
18   podSelector:
19     matchLabels:
20       app: node-app
21   ingress:
22     - ports:
23       - port: 8888

```

Listing 5: Network Policy handling traffic between Pod and external components

The other scenarios cannot be enforced with NetworkPolicies, since the rules depend on individual containers instead of Pods. The rules need to be applied inside the Pod's network namespace with tools such as IPTables.

5.1.1 IPTables

Usually, IPTables rules are applied to packets by their IP addresses and ports. However, modern versions of IPTables ship with an owner module, which supports rules based on application's user, group, and session identifiers. For example, `iptables -A OUTPUT -m owner --uid-owner 1000 -j REJECT` would reject all egress packets from containers with the user id. While the session identifier is assigned at runtime

by the kernel, containers' user and group ids can be overridden with Pod definition's `runAsUser` and `runAsGroup` fields, respectively. Thus, any ingress or egress traffic of a container with unique user identifier can be filtered with IPTables.

Listing 6 shows IPTables rules that can enforce network rules in other scenarios. As ingress rules, the main application port is opened on all network devices, while sidecar port is only allowed on loopback device. For egress traffic, the main application is given access to both containers' ports, while the sidecar can only access its own port. All other traffic is rejected by default.

```
1 #!/bin/bash
2 APP_UID=1000
3 APP_PORT=8888
4 SIDECAR_UID=2000
5 SIDECAR_PORT=8125
6
7 # Ingress rules
8 iptables -A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j
  ACCEPT
9 iptables -A INPUT -p tcp --dport $APP_PORT -j ACCEPT # Scenarios 1,
  2
10 iptables -A INPUT -i lo -p tcp --dport $SIDECAR_PORT -j ACCEPT #
  Scenario 3
11 iptables -A INPUT -j DROP
12
13 # Egress rules
14 iptables -A OUTPUT -m conntrack --ctstate RELATED,ESTABLISHED -j
  ACCEPT
15
16 ## Main app rules
17 iptables -A OUTPUT -d 127.0.0.1 -m owner --uid-owner $APP_UID -p
  tcp --dport $APP_PORT -j ACCEPT # Scenario 2
18 iptables -A OUTPUT -d 127.0.0.1 -m owner --uid-owner $APP_UID -p
  tcp --dport $SIDECAR_PORT -j ACCEPT # Scenario 2
19
20 ## Sidecar rules
21 iptables -A OUTPUT -d 127.0.0.1 -m owner --uid-owner $SIDECAR_UID -
  p tcp --dport $SIDECAR_PORT -j ACCEPT
22
23 iptables -A OUTPUT -j REJECT
```

Listing 6: IPTables rules for the Pod

The approach is similar to how Istio redirects all Pod traffic through the Envoy sidecar and the service mesh. Istio adds IPTables rules that reroute all traffic, except that which originates from user id 1337, to the Envoy proxy [48]. The user id itself is reserved for the proxy, and it is used to avoid re-routing the proxy traffic, which would cause an infinite loop.

eBPF programs were also considered as an option for IPTables. Since XDP programs support only ingress rules, they are not a valid option for solving the problem. However, a TC program that could distinguish sidecar container traffic from others could be a possible option. TC programs manipulate packets by modifying struct `__sk_buff` in the user space. The complete C structure is listed as Example

7.

```
1 /* user accessible mirror of in-kernel sk_buff.
2  * new fields can only be added to the end of this structure
3  */
4 struct __sk_buff {
5     __u32 len;
6     __u32 pkt_type;
7     __u32 mark;
8     __u32 queue_mapping;
9     __u32 protocol;
10    __u32 vlan_present;
11    __u32 vlan_tci;
12    __u32 vlan_proto;
13    __u32 priority;
14    __u32 ingress_ifindex;
15    __u32 ifindex;
16    __u32 tc_index;
17    __u32 cb[5];
18    __u32 hash;
19    __u32 tc_classid;
20    __u32 data;
21    __u32 data_end;
22    __u32 napi_id;
23
24    /* Accessed by BPF_PROG_TYPE_sk_skb types from here to ... */
25    __u32 family;
26    __u32 remote_ip4;    /* Stored in network byte order */
27    __u32 local_ip4;     /* Stored in network byte order */
28    __u32 remote_ip6[4]; /* Stored in network byte order */
29    __u32 local_ip6[4];  /* Stored in network byte order */
30    __u32 remote_port;   /* Stored in network byte order */
31    __u32 local_port;    /* stored in host byte order */
32    /* ... here. */
33
34    __u32 data_meta;
35    __bpf_md_ptr(struct bpf_flow_keys *, flow_keys);
36    __u64 tstamp;
37    __u32 wire_len;
38    __u32 gso_segs;
39    __bpf_md_ptr(struct bpf_sock *, sk);
40    __u32 gso_size;
41    __u8  tstamp_type;
42    __u32 :24;    /* Padding, future use. */
43    __u64 hwtstamp;
44 };
```

Listing 7: skbuff struct in Linux v6.3 [31]

As can be seen in the example, the structure has no identifiers that are unique for each container. A notable attribute in the structure is the mark field, which can be used as an identifier for the packet in the kernel. It does not propagate with the packet, which means that marking can only be used for traffic filtering before the packet leaves the network stack. If combined with iptables rule that

uniquely marks the packets egressing from the sidecar, for example by setting the mark as 2 `iptables -t mangle -A PREROUTING -m owner --uid-owner 1000 -j MARK --set-mark 2`, a TC eBPF program could be used for filtering the traffic. However, using eBPF brings an extra layer of complexity in comparison to simply using IPTables for traffic filtering. Furthermore, eBPF programs do not provide any performance gains for egressing traffic, since the TC hooks execute after the packet has been processed by IPTables.

5.1.2 Deployment

The commands for creating the IPTables rules could be applied to the Pod with the `initContainer` or `postStart` lifecycle handler. However, the commands require root user permissions with `NET_ADMIN` and `NET_RAW` capabilities. Using these methods would require changing Pod's security admission level to privileged. Thus, it is more optimal to apply the rules outside the Pod's context; for example, by using a `DaemonSet`, similarly to how many CNI plugins configure the Pod network.

5.2 Own network namespace for sidecar

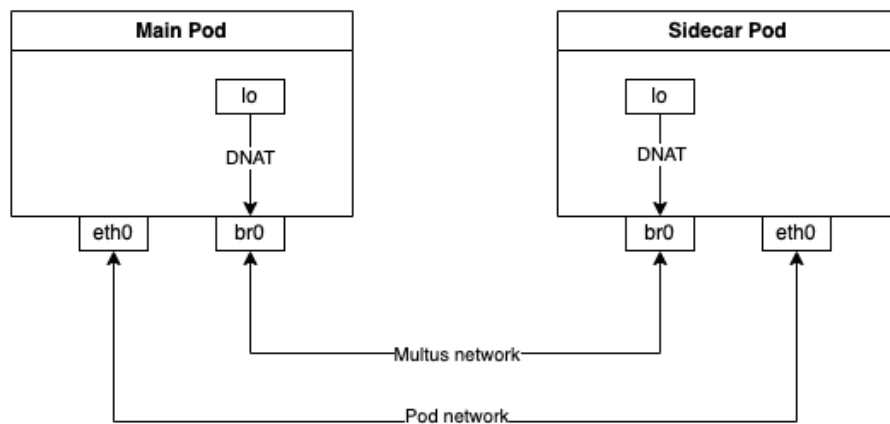


Figure 5: Multiple network namespaces architecture

The figure 5 illustrates the network architecture for the solution where containers are deployed in different Pods. The idea behind this approach is to use Network Policies for filtering the traffic between the containers. Since the containers are not technically sidecars anymore, the solution needs to recreate communication via localhost address.

Although it is possible to use the default network created by the CNI plugin for sidecar traffic, the proposed solution creates a new isolated network for traffic with Multus. The isolation allows the use of separate IP address pool, which comes in handy when re-creating the communication via localhost.

5.2.1 Creating isolated network for sidecar traffic

It is completely possible to use the default Pod network created by the CNI plugin for traffic between sidecars. With this setup, the network rules are simply Network Policies on top of the policy that implicitly denies all other traffic.

It is also possible to build new network interfaces for sidecar traffic with Multus.

```
1 apiVersion: k8s.cni.cncf.io/v1
2 kind: NetworkAttachmentDefinition
3 metadata:
4   name: localhost-replacement
5   namespace: app
6 spec:
7   config: '{
8     "cniVersion": "0.3.1",
9     "name": "localhost-replacement",
10    "capabilities": { "ips": true },
11    "type": "macvlan",
12    "master": "eth0",
13    "mode": "bridge",
14    "ipam": {
15      "type": "static"
16    }
17  }'
```

Listing 8: NetworkAttachmentDefinition for creating MacVLAN bridge between the Pods

The setup requires creating NetworkAttachmentDefinition CRD for the network. Each Pod that should be part of the network needs a `k8s.v1.cni.cncf.io/networks` annotation, as seen in example 9.

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: node-app
5   namespace: app
6 spec:
7   replicas: 1
8   selector:
9     matchLabels:
10      app: node-app
11   template:
12     metadata:
13       labels:
14         app: node-app
15     annotations:
16       k8s.v1.cni.cncf.io/networks: '[{'
```

```

17         "name": "localhost-replacement",
18         "interface": "br0",
19         "ips": [ "192.168.1.201/24" ]
20     }]'
21     spec:
22         containers:
23         - name: node-app
24           image: node-app
25 ---
26 apiVersion: apps/v1
27 kind: Deployment
28 metadata:
29     name: node-app-statsd
30     namespace: app
31 spec:
32     replicas: 1
33     selector:
34         matchLabels:
35             app: node-app-statsd
36     template:
37         metadata:
38             labels:
39                 app: node-app-statsd
40         annotations:
41             k8s.v1.cni.cncf.io/networks: '[{
42                 "name": "localhost-replacement",
43                 "interface": "br0",
44                 "ips": [ "192.168.1.202/24" ]
45             }]'
46     spec:
47         containers:
48         - name: statsd
49           image: hypnza/statsd_dumpmessages

```

Listing 9: Deployment with Pod connected to the bridge

The common NetworkPolicies are not applied to the Multus network. However, the Multus developers have introduced another CRD, MultiNetworkPolicy [38], which mimics NetworkPolicies but only works for networks created by Multus. The resource is identical to NetworkPolicies, as seen in example 10.

```

1 apiVersion: k8s.cni.cncf.io/v1beta1
2 kind: MultiNetworkPolicy
3 metadata:
4     name: deny-all-br0
5     namespace: app
6     annotations:

```

```

7     k8s.v1.cni.cncf.io/policy-for: localhost-replacement
8 spec:
9   podSelector: {}
10  policyTypes:
11    - Ingress
12    - Egress
13 ---
14 apiVersion: k8s.cni.cncf.io/v1beta1
15 kind: MultiNetworkPolicy
16 metadata:
17   name: allow-statsd
18   namespace: app
19   annotations:
20     k8s.v1.cni.cncf.io/policy-for: localhost-replacement
21 spec:
22   podSelector:
23     matchLabels:
24       app: node-app-statsd
25   ingress:
26     - from:
27       - podSelector:
28         matchLabels:
29           app: node-app
30     ports:
31       - protocol: TCP
32         port: 8125

```

Listing 10: Zero trust MultiNetworkPolicies for br0 network interface

However, the Multus CNI does not enforce the MultiNetworkPolicy rules in any way. The implementation is left for other DaemonSets like multi-networkpolicy-iptables [39] and multi-networkpolicy-tc [40], both of which are created by the Multus team. As the names suggest, the DaemonSets implement the enforcement with IPTables and eBPF TC programs, respectively. The multi-networkpolicy-iptables is used for the solution, since iptables implements all other network rules for the cluster.

With these configurations and dependencies, the containers should have a new network and reach each other via the bridge interface. As a next step, the solution needs to find a way for re-routing the traffic destined for localhost to the new network. The script in listing 11 shows how to redirect all egressing traffic with IPTables from the source of 127.0.0.1:8125 to a network interface with external IP address, like 192.168.1.202:8125. By default, the re-routing of local traffic outside the local machine is denied by the kernel. To overrule this default, the script needs to set a flag `net.ipv4.conf.all.route_localnet=1`.

```

1  #!/bin/bash
2  sysctl -w net.ipv4.conf.all.route_localnet=1

```

```
3 iptables -t nat -A OUTPUT -p ALL --dport 8125 -j DNAT --to-  
destination 192.168.1.202:8125  
4 iptables -t nat -A POSTROUTING -j MASQUERADE
```

Listing 11: Script for mapping localhost to br0 interface

5.2.2 Mutating sidecars to multiple Pods?

MutatingAdmissionWebhook?

6 Solution Evaluation

7 Discussion

8 Conclusion

References

- [1] AppArmor. *AppArmor*. 2022. URL: <https://apparmor.net/> (visited on 03/31/2023).
- [2] Docker Authors. *Docker overview*. 2023. URL: <https://docs.docker.com/get-started/overview/> (visited on 04/06/2023).
- [3] Docker Authors. *Use containers to Build, Share and Run your applications*. 2023. URL: <https://www.docker.com/resources/what-container/> (visited on 04/06/2023).
- [4] etcd Authors. *etcd*. 2023. URL: <https://etcd.io/> (visited on 04/05/2023).
- [5] Kubernetes Authors. *Admission Controllers Reference*. 2023. URL: <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/> (visited on 04/06/2023).
- [6] Kubernetes Authors. *Cluster Networking*. 2022. URL: <https://kubernetes.io/docs/concepts/cluster-administration/networking/> (visited on 03/09/2023).
- [7] Kubernetes Authors. *Deployments*. 2023. URL: <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/> (visited on 04/04/2023).
- [8] Kubernetes Authors. *Kubernetes*. 2023. URL: <https://kubernetes.io/> (visited on 03/31/2023).
- [9] Kubernetes Authors. *Kubernetes components*. 2023. URL: <https://kubernetes.io/docs/concepts/overview/components/> (visited on 04/05/2023).
- [10] Kubernetes Authors. *Pod Security Admission*. 2023. URL: <https://kubernetes.io/docs/concepts/security/pod-security-admission/> (visited on 04/26/2023).
- [11] Kubernetes Authors. *Pod Security Standards*. 2023. URL: <https://kubernetes.io/docs/concepts/security/pod-security-standards/> (visited on 04/04/2023).
- [12] Kubernetes Authors. *Pods*. 2023. URL: <https://kubernetes.io/concepts/workloads/pods/> (visited on 04/04/2023).
- [13] Kubernetes Authors. *Resource Management for Pods and Containers*. 2023. URL: <https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/> (visited on 05/30/2023).
- [14] Kubernetes Authors. *Services*. 2023. URL: <https://kubernetes.io/docs/concepts/services-networking/service/> (visited on 04/04/2023).
- [15] The CNI Authors. *Container Network Interface (CNI) Specification*. 2023. URL: <https://www.cni.dev/docs/spec/> (visited on 04/20/2023).

- [16] Gerald Budigiri et al. “Network policies in kubernetes: Performance evaluation and security analysis”. In: *2021 Joint European Conference on Networks and Communications & 6G Summit (EuCNC/6G Summit)*. IEEE. 2021, pp. 407–412.
- [17] Thanh Bui. “Analysis of docker security”. In: *arXiv preprint arXiv:1501.02967* (2015).
- [18] Brendan Burns and David Oppenheimer. “Design patterns for container-based distributed systems”. In: *8th {USENIX} workshop on hot topics in cloud computing (HotCloud 16)*. 2016.
- [19] Brendan Burns et al. “Borg, Omega, and Kubernetes”. In: *Communications of the ACM* 59.5 (2016), pp. 50–57.
- [20] Theo Combe, Antony Martin, and Roberto Di Pietro. “To docker or not to docker: A security perspective”. In: *IEEE Cloud Computing* 3.5 (2016), pp. 54–62.
- [21] Cilium Developers. *Cilium*. 2023. URL: <https://cilium.io/> (visited on 03/14/2023).
- [22] Cilium Developers. *Component overview*. 2023. URL: <https://docs.cilium.io/en/v1.13/overview/component-overview/> (visited on 04/12/2023).
- [23] Cilium Developers. *Component overview*. 2023. URL: <https://docs.cilium.io/en/v1.13/security/policy/language/> (visited on 04/12/2023).
- [24] Cilium Developers. *Kubernetes Without kube-proxy*. 2023. URL: <https://docs.cilium.io/en/v1.13/network/kubernetes/kubeproxy-free> (visited on 04/12/2023).
- [25] Cilium Developers. *Transparent Encryption*. 2023. URL: <https://docs.cilium.io/en/v1.13/security/network/encryption/> (visited on 04/12/2023).
- [26] Cilium developers. *Program types*. 2018. URL: <https://docs.cilium.io/en/latest/bpf/progtypes/> (visited on 03/16/2023).
- [27] Istio developers. *The Istio service mesh*. 2023. URL: <https://istio.io/latest/about/service-mesh/> (visited on 04/06/2023).
- [28] Linux Developers. *capabilities(7) — Linux manual page*. 2023. URL: <https://man7.org/linux/man-pages/man7/capabilities.7.html> (visited on 03/31/2023).
- [29] Linux Developers. *namespaces(7) — Linux manual page*. 2023. URL: <https://man7.org/linux/man-pages/man7/namespaces.7.html> (visited on 03/31/2023).
- [30] Linux Developers. *seccomp(2) — Linux manual page*. 2023. URL: <https://man7.org/linux/man-pages/man2/seccomp.2.html> (visited on 03/31/2023).

- [31] Linux developers. *bpf.h in Linux source code*. 2023. URL: <https://github.com/torvalds/linux/blob/v6.3/include/uapi/linux/bpf.h> (visited on 06/23/2023).
- [32] Docker. *Docker default Seccomp profile*. 2022. URL: <https://github.com/moby/moby/blob/23.0/profiles/seccomp/default.json> (visited on 03/31/2023).
- [33] Docker. *Docker overview*. 2018. URL: <https://docs.docker.com/get-started/overview/> (visited on 03/16/2023).
- [34] Docker. *Docker security*. 2023. URL: <https://docs.docker.com/engine/security/> (visited on 03/31/2023).
- [35] Docker. *libnetwork*. 2023. URL: <https://github.com/moby/moby/tree/master/libnetwork> (visited on 03/10/2023).
- [36] Martin Fowler and James Lewis. *Microservices*. 2014. URL: <https://martinfowler.com/articles/microservices.html> (visited on 05/09/2023).
- [37] Thomas Graf. *Cilium Service Mesh – Everything You Need to Know*. 2023. URL: <https://isovalent.com/blog/post/cilium-service-mesh/> (visited on 04/12/2023).
- [38] Kubernetes Network Plumbing Working Group. *multi-networkpolicy CRD*. 2023. URL: <https://github.com/k8snetworkplumbingwg/multi-networkpolicy> (visited on 07/03/2023).
- [39] Kubernetes Network Plumbing Working Group. *multi-networkpolicy-iptables*. 2023. URL: <https://github.com/k8snetworkplumbingwg/multi-networkpolicy-iptables> (visited on 07/03/2023).
- [40] Kubernetes Network Plumbing Working Group. *multi-networkpolicy-tc*. 2023. URL: <https://github.com/k8snetworkplumbingwg/multi-networkpolicy-tc> (visited on 07/03/2023).
- [41] Network Plumbing Working Group. *Multus-CNI*. 2023. URL: <https://github.com/k8snetworkplumbingwg/multus-cni> (visited on 04/25/2023).
- [42] Network Plumbing Working Group. *Multus-CNI Configuration Reference*. 2023. URL: <https://github.com/k8snetworkplumbingwg/multus-cni/blob/master/docs/configuration.md> (visited on 04/25/2023).
- [43] The Kubernetes Networking Guide. *Calico*. 2023. URL: <https://www.tknng.io/cni/calico/> (visited on 03/14/2023).
- [44] The Kubernetes Networking Guide. *Cilium*. 2023. URL: <https://www.tknng.io/cni/cilium/> (visited on 03/14/2023).
- [45] The Kubernetes Networking Guide. *CNI*. 2023. URL: <https://www.tknng.io/cni/> (visited on 04/20/2023).
- [46] Michael Hausenblas. *Container Networking*. O'Reilly Media, Incorporated, 2018.

- [47] Toke Høiland-Jørgensen et al. “The express data path: Fast programmable packet processing in the operating system kernel”. In: *Proceedings of the 14th international conference on emerging networking experiments and technologies*. 2018, pp. 54–66.
- [48] Istio. *Install Istio with the Istio CNI plugin*. 2023. URL: <https://istio.io/v1.17/docs/setup/additional-setup/cni/> (visited on 06/14/2023).
- [49] Pooyan Jamshidi et al. “Microservices: The journey so far and challenges ahead”. In: *IEEE Software* 35.3 (2018), pp. 24–35.
- [50] Alper Kerman et al. “Implementing a zero trust architecture”. In: *National Institute of Standards and Technology (NIST)* (2020).
- [51] Xin Lin et al. “A measurement study on linux container security: Attacks and countermeasures”. In: *Proceedings of the 34th Annual Computer Security Applications Conference*. 2018, pp. 418–429.
- [52] Steven McCanne and Van Jacobson. “The BSD Packet Filter: A New Architecture for User-level Packet Capture.” In: *USENIX winter*. Vol. 46. 1993.
- [53] Dirk Merkel et al. “Docker: lightweight linux containers for consistent development and deployment”. In: *Linux j* 239.2 (2014), p. 2.
- [54] Sebastiano Miano et al. “A framework for eBPF-based network functions in an era of microservices”. In: *IEEE Transactions on Network and Service Management* 18.1 (2021), pp. 133–151.
- [55] Microsoft. *Threat Matrix for Kubernetes*. 2023. URL: <https://microsoft.github.io/Threat-Matrix-for-Kubernetes/> (visited on 05/24/2023).
- [56] Shixiong Qi, Sameer G Kulkarni, and KK Ramakrishnan. “Assessing container network interface plugins: Functionality, performance, and scalability”. In: *IEEE Transactions on Network and Service Management* 18.1 (2020), pp. 656–671.
- [57] Scott Rose et al. *Zero trust architecture*. Tech. rep. National Institute of Standards and Technology, 2020.
- [58] Nataliya Shevchenko et al. *Threat modeling: a summary of available methods*. Tech. rep. Carnegie Mellon University Software Engineering Institute Pittsburgh United . . . , 2018.
- [59] Tigera. *About Calico*. 2023. URL: <https://docs.tigera.io/calico/3.25/about> (visited on 04/13/2023).
- [60] Tigera. *About Network Policy*. 2023. URL: <https://docs.tigera.io/calico/latest/about/about-network-policy> (visited on 04/13/2023).
- [61] Tigera. *Install calicoctl*. 2023. URL: <https://docs.tigera.io/calico/3.25/operations/calicoctl/install> (visited on 04/13/2023).
- [62] Marcos AM Vieira et al. “Fast packet processing with ebpf and xdp: Concepts, code, challenges, and applications”. In: *ACM Computing Surveys (CSUR)* 53.1 (2020), pp. 1–36.

A Dockerfile for penetration testing

```
1 FROM --platform=linux/amd64 ubuntu:22.04
2
3 RUN apt update
4 RUN apt install -y curl
5
6 # Setup kubectl
7 RUN curl -fsSLo /etc/apt/keyrings/kubernetes-archive-keyring.gpg
   https://packages.cloud.google.com/apt/doc/apt-key.gpg
8 RUN echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-archive-
   keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main'
   | tee /etc/apt/sources.list.d/kubernetes.list
9
10 # Install deps
11 RUN apt update
12 RUN apt install -y net-tools nmap ncat kubectl etcd iputils-ping
   iproute2
13
14 # kubeletctl
15 RUN curl -LO https://github.com/cyberark/kubeletctl/releases/
   download/v1.9/kubeletctl_linux_amd64 && chmod a+x ./
   kubeletctl_linux_amd64 && mv ./kubeletctl_linux_amd64 /usr/local
   /bin/kubeletctl
16
17 RUN apt install -y iptables
18
19 # Point to the internal API server hostname
20 RUN echo 'export APISERVER=https://kubernetes.default.svc' >> root
   /.bashrc
21 # Path to ServiceAccount token
22 RUN echo 'export SERVICEACCOUNT=/var/run/secrets/kubernetes.io/
   serviceaccount' >> root/.bashrc
23 # Read this Pod's namespace
24 RUN echo 'export NAMESPACE="$(cat ${SERVICEACCOUNT}/namespace)"' >>
   root/.bashrc
25 # Read the ServiceAccount bearer token
26 RUN echo 'export TOKEN="$(cat ${SERVICEACCOUNT}/token)"' >> root/.
   bashrc
27 # Reference the internal certificate authority (CA)
28 RUN echo 'export CACERT="${SERVICEACCOUNT}/ca.crt"' >> root/.bashrc
```

Listing 12: Dockerfile for penetration testing

B Example webserver deployed with StatsD sidecar

```
1 const Koa = require('koa');
2 const app = new Koa();
3
4 const StatsD = require('node-statsd');
5 const client = new StatsD();
6
7 const logMessage = async (ctx, next) => {
8   await next();
9   const rt = ctx.response.get('X-Response-Time');
10  console.log(`${ctx.method} ${ctx.url} - ${rt}`);
11  client.timing('response_time', rt);
12  client.increment('response_counter');
13 }
14
15 const setResponseTimeCtx = async (ctx, next) => {
16   const start = Date.now();
17   await next();
18   const ms = Date.now() - start;
19   ctx.set('X-Response-Time', `${ms}ms`);
20 }
21
22 app.use(logMessage);
23 app.use(setResponseTimeCtx);
24 app.use(async ctx => {
25   ctx.body = 'Hello World';
26 });
27
28 app.listen(8888);
```

Listing 13: Node.js server (index.js)

```

1 {
2   "name": "node-app",
3   "version": "1.0.0",
4   "description": "Webserver w/ StatsD",
5   "main": "index.js",
6   "scripts": {
7     "start": "node index.js"
8   },
9   "author": "Aarni Halinen",
10  "license": "ISC",
11  "dependencies": {
12    "koa": "^2.14.2",
13    "node-statsd": "^0.1.1"
14  }
15 }

```

Listing 14: package.json

```

1 FROM node:18-alpine
2
3 WORKDIR /app
4
5 COPY index.js package.json package-lock.json ./
6 RUN npm install --omit=dev
7
8 ENTRYPOINT [ "npm", "start" ]

```

Listing 15: Dockerfile for the server

C Webserver and StatsD deployed with Multus networking

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: node-app
5   namespace: app
6 spec:
7   replicas: 1
8   selector:
9     matchLabels:
10      app: node-app
11   template:
12     metadata:
13       labels:
14         app: node-app
15     annotations:
16       k8s.v1.cni.cncf.io/networks: '[{
17         "name": "localhost-replacement",
18         "interface": "br0",
19         "ips": [ "192.168.1.201/24" ]
20       }]'
21   spec:
22     affinity:
23       nodeAffinity:
24         requiredDuringSchedulingIgnoredDuringExecution:
25           nodeSelectorTerms:
26             - matchExpressions:
27               - key: kubernetes.io/hostname
28                 operator: In
29                 values:
30                   - minikube-m02
31     automountServiceAccountToken: false
32   containers:
33     - name: node-app
34       image: node-app
35       command: [ "node" ]
36       args: [ "index.js" ]
37       ports:
38         - containerPort: 8888
39           protocol: TCP
40       securityContext:
```

```

41         allowPrivilegeEscalation: true
42         readOnlyRootFilesystem: false
43         runAsUser: 0
44         runAsGroup: 0
45         runAsNonRoot: false
46         seccompProfile:
47             type: RuntimeDefault
48         capabilities:
49             drop:
50                 - ALL
51     resources:
52         limits:
53             cpu: "1"
54             memory: 1024Mi
55 ---
56 apiVersion: v1
57 kind: Service
58 metadata:
59     name: node-app
60     namespace: app
61 spec:
62     ports:
63     - port: 8888
64       protocol: TCP
65       nodePort: 30001
66     selector:
67         app: node-app
68     type: NodePort

```

Listing 16: Server's K8s deployment resource


```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: node-app-statsd
5   namespace: app
6 spec:
7   replicas: 1
8   selector:
9     matchLabels:
10      app: node-app-statsd
11 template:
12   metadata:
13     labels:
14       app: node-app-statsd
15     annotations:
16       k8s.v1.cni.cncf.io/networks: '[{
17         "name": "localhost-replacement",
18         "interface": "br0",
19         "ips": [ "192.168.1.202/24" ]
20       }]'
21   spec:
22     affinity:
23       nodeAffinity:
24         requiredDuringSchedulingIgnoredDuringExecution
25       :
26         nodeSelectorTerms:
27         - matchExpressions:
28           - key: kubernetes.io/hostname
29             operator: In
30             values:
31               - minikube-m02
32         automountServiceAccountToken: false
33     containers:
34     - name: statsd
35       # image: statsd/statsd
36       image: hypnza/statsd_dumpmessages
37       ports:
38       - containerPort: 8125
39       securityContext:
40         allowPrivilegeEscalation: false
41         readOnlyRootFilesystem: true
42         runAsUser: 101
43         runAsGroup: 101
44         runAsNonRoot: true
45         seccompProfile:

```

```
45         type: RuntimeDefault
46     capabilities:
47         drop:
48             - ALL
49     resources:
50         limits:
51             cpu: "1"
52             memory: 256Mi
```

Listing 17: StatsD container's K8s Deployment

```

1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: deny-all
5   namespace: app
6 spec:
7   podSelector: {}
8   policyTypes:
9     - Ingress
10    - Egress
11 ---
12 apiVersion: k8s.cni.cncf.io/v1
13 kind: NetworkAttachmentDefinition
14 metadata:
15   name: localhost-replacement
16   namespace: app
17 spec:
18   config: '{
19     "cniVersion": "0.3.1",
20     "name": "localhost-replacement",
21     "capabilities": { "ips": true },
22     "type": "macvlan",
23     "master": "eth0",
24     "mode": "bridge",
25     "ipam": {
26       "type": "static"
27     }
28   }',
29 ---
30 apiVersion: k8s.cni.cncf.io/v1beta1
31 kind: MultiNetworkPolicy
32 metadata:
33   name: deny-all-br0
34   namespace: app
35   annotations:
36     k8s.v1.cni.cncf.io/policy-for: localhost-replacement
37 spec:
38   podSelector: {}
39   policyTypes:
40     - Ingress
41     - Egress
42 ---
43 apiVersion: k8s.cni.cncf.io/v1beta1
44 kind: MultiNetworkPolicy
45 metadata:

```

```

46  name: allow-statsd
47  namespace: app
48  annotations:
49    k8s.v1.cni.cncf.io/policy-for: localhost-replacement
50 spec:
51   podSelector:
52     matchLabels:
53       app: node-app-statsd
54   ingress:
55     - from:
56         - podSelector:
57             matchLabels:
58               app: node-app
59       ports:
60         - protocol: TCP
61           port: 8125

```

Listing 18: Network interfaces and policies